

```

%matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

```

## Univariate Linear regression

```

In [19]: data=pd.read_csv("C:/Users/Soham/ML/Labs/dataset1.txt", header=None)
data.head()

Out[19]:
   0      1
0  6.1101  17.5920
1  5.5277   9.1302
2  8.5186  13.6620
3  7.0032  11.8540
4  5.8598   6.8233

In [20]: data.describe()


Out[20]:
   0      1
count  97.000000  97.000000
mean    8.159800   5.839135
std     3.869884   5.510262
min     5.026900  -2.680700
25%    5.707700   1.986900
50%    6.589400   4.562300
75%    8.578100   7.046700
max    22.203000  24.147000

In [21]: data.columns = ['Population','Profit']

In [22]: plt.scatter(data['Population'],data['Profit'])
plt.xticks(np.arange(5,30,step=5))
plt.yticks(np.arange(-5,30,step=5))
plt.xlabel('Population (in 10,000s)')
plt.ylabel('Profit (in 10,000$)')
plt.title('Profit vs Population')

Out[22]: Text(0.5, 1.0, 'Profit vs Population')

```



## Cost function $J(\theta)$

```

In [23]: def computeCost(X,y,theta):
        """
        Take in a numpy array X,y,theta and get cost function using theta as parameter in a linear regression model
        """
        m=len(y)
        prediction =X.dot(theta)
        square_err = (prediction -y)**2

        return 1/(2*m)*np.sum(square_err)

In [24]: data['x0'] =1

In [25]: data_val= data.values
m = len(data_val[:-1])
X =data[['x0','Population']].iloc[:-1].values
y = data['Profit'].iloc[:-1].values.reshape(m,1)
theta = np.zeros((2,1))

m, X.shape, y.shape, theta.shape

Out[25]: (96, (96, 2), (96, 1), (2, 1))

h(θ) = x0θ0 +x1θ1 ....(x0 = 1)

In [26]: computeCost(X,y,theta)

Out[26]: 32.40484177877831

In [27]: data.tail()

Out[27]:
   Population  Profit  x0
92      5.8707  7.20290   1
93      5.3054  1.98690   1
94      8.2934  0.14454   1
95     13.3940  9.05510   1
96      5.4369  0.61705   1

```

## Gradient Descent

```

In [28]: def gradientDescent(X,y,theta,alpha,num_iters):
        """
        Take numpy array for X,y,theta and update theta for every iteration of gradient steps
        return theta adn the list of cost of theta during each iteration
        """
        m=len(y)
        J_history=[]
        for i in range(num_iters):
            predictions= X.dot(theta)
            error =np.dot(X.transpose()),(predictions - y))
            descent= alpha * 1/m *error
            theta-= descent
            J_history.append(computeCost(X,y,theta))

        return theta,J_history


In [29]: theta, J_history = gradientDescent(X,y,theta,0.001,2000)

In [30]: print(f" h(x) = {str(round(theta[0,0],2))} + {str(round(theta[1,0],2))}x1")
h(x) = -1.11 + 0.92x1

In [31]: from mpl_toolkits.mplot3d import Axes3D
#Generating values for theta0, theta1 and the resulting cost value
theta0_vals=np.linspace(-10,10,100)
theta1_vals=np.linspace(-1,4,100)
J_vals=np.zeros((len(theta0_vals),len(theta1_vals)))
for i in range(len(theta0_vals)):
    for j in range(len(theta1_vals)):
        t=np.array([theta0_vals[i],theta1_vals[j]])
        J_vals[i,j]=computeCost(X,y,t)

#Generating the surface plot
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
surf=ax.plot_surface(theta0_vals,theta1_vals,J_vals,cmap="coolwarm")
fig.colorbar(surf, shrink=0.5, aspect=5)
ax.set_xlabel("$\theta_0$")
ax.set_ylabel("$\theta_1$")
ax.set_zlabel("$J(\theta)$")
#rotate for better angle
ax.view_init(30,120)

```



```

In [32]: plt.plot(J_history)
plt.xlabel("Iteration")
plt.ylabel("$J(\theta)$")
plt.title("Cost function using Gradient Descent")

Out[32]: Text(0.5, 1.0, 'Cost function using Gradient Descent')

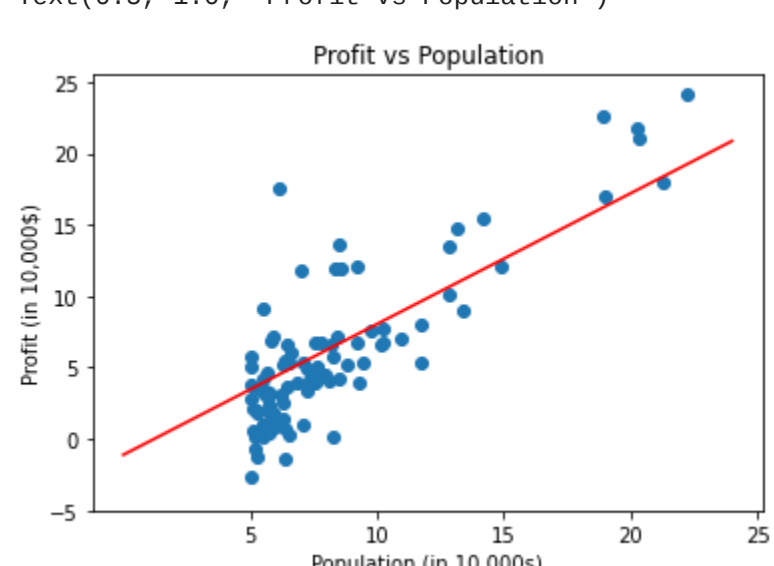
Cost function using Gradient Descent

Iteration
J(theta)

In [33]: plt.scatter(data['Population'], data['Profit'])
x_value = [x for x in range(25)]
y_value = [x*theta[1] + theta[0] for x in x_value]
plt.plot(x_value, y_value, color = 'r')
plt.xticks(np.arange(5,30,step=5))
plt.yticks(np.arange(-5,30,step=5))
plt.xlabel('Population (in 10,000s)')
plt.ylabel('Profit (in 10,000$)')
plt.title('Profit vs Population')

Out[33]: Text(0.5, 1.0, 'Profit vs Population')

```



```

In [34]: def predict(x,theta):
        """
        takes in numpy array x and theta and returns predicted value of y
        """
        predictions = np.dot(theta.transpose(),x)
        return predictions[0]

In [35]: data.tail(1)

Out[35]:
   Population  Profit  x0
96      5.4369  0.61705   1

Multivariate Linear Regression

Imports

```

```

In [36]: import statsmodels.api as sm
from sklearn.linear_model import LinearRegression
np.random.seed(123)

In [37]: data = pd.read_csv("C:/Users/Soham/ML/Labs/dataset2.csv")
data.head()

Out[37]:
   Size of the house (in square feet)  Number of bedrooms  Price of the house
0                                2104                      3          399900
1                                1600                      3          329900
2                                2400                      3          369000
3                                1416                      2          232000
4                                3000                      4          539900

In [38]: data.isnull().sum()

Out[38]:
Size of the house (in square feet)    0
Number of bedrooms                   0
Price of the house                   0
dtype: int64

In [39]: def normalize(dataframe):
        df = dataframe.copy()
        for col in df.columns:
            df[col] = (df[col]-df[col].mean())/df[col].std()
        return df

In [40]: normalized_data = normalize(data)
normalized_data.head()

Out[40]:
   Size of the house (in square feet)  Number of bedrooms  Price of the house
0                                0.130010             -0.223675             0.475747
1                               -0.504190             -0.223675             -0.084074
2                                0.502476             -0.223675             0.228626
3                               -0.735723             -1.537767             -0.867025
4                                1.257476              1.090417             1.595389

```

```

In [41]: X = normalized_data.iloc[:, :-1].values
y = normalized_data.iloc[:, -1].values

In [42]: m = y.size
n = data.shape[1]

In [43]: y.shape

Out[43]: (47, )

In [44]: y = y.reshape(m,1)
y.shape

Out[44]: (47, 1)

In [45]: ones = np.ones((m,1))
X1 = np.concatenate((ones,X),axis=1)
X1[:5]

Out[45]: array([[ 1.          ,  0.13000987, -0.22367519],
 [ 1.          , -0.50418984, -0.22367519],
 [ 1.          ,  0.50247636, -0.22367519],
 [ 1.          , -0.73572308, -1.53776691],
 [ 1.          ,  1.25747602,  1.09041654]])

In [46]: alpha = 0.01
theta = np.random.rand(n,1)
epoch = 10000

```

```

In [47]: def GD(X1,y,theta,alpha,decimals=5):
        past_Cost = []
        past_theta = [theta]
        m = y.size
        n = X1.shape[1]
        for i in range(epoch):
            h_theta = np.dot(X1,theta)
            error = h_theta-y
            cost = np.dot(error.T, error)/(2*m)
            past_cost.append(cost[0][0])
            diff = np.dot(X1.T, error)/m
            theta = theta - (alpha*diff)
            past_theta.append(theta)
            # Task 4 :- do early stopping (I have considered 5 decimal places, you can change the decimals parameter if you want)
            if np.equal(np.round(past_theta[i],decimals=decimals),np.round(past_theta[i+1],decimals=decimals)).sum() == n:
                break
        return past_cost, past_theta, i+1


In [48]: pastCost, pastTheta,stop_epoch = GD(X1=X1, y=y, theta=theta, epoch=epoch,alpha=alpha)

In [49]: print(f" Our model performed {stop_epoch} epochs out of {epoch} epochs before converging'")
Our model performed 1320 epochs out of 10000 epochs before converging

In [50]: plt.plot(pastCost)

Out[50]: <matplotlib.lines.Line2D at 0x1c7bf4d3d90>

```



```

In [51]: best_theta = np.array(pastTheta[-1]).reshape(n,)
print(best_theta)

[ 1.20603184e-06  8.83291779e-01 -5.17046112e-02]

In [52]: print(f"Parameters from StatsModels -> {sm.OLS(y,X1).fit().params}")
print(f"Parameters from ScikitLearn -> {LinearRegression().fit(X1,y).coef_}")

Parameters from StatsModels -> [-9.71445147e-17  8.84765988e-01 -5.31788197e-02]
Parameters from ScikitLearn -> [[ 0.          0.88476599 -0.05317882]]

In [ ]:

```